

# A Simple Analysis for *Divorce* Model

Jingyue Wu

Department of Computing and Information Science, Cornell University

Advisor: Yang Ning, Department of Statistics and Data Science

2019, December 10

## 1 Executive Summary

The goal of this analysis is to create models that could accurately predict the divorce based on all the available features, and also find other significant items and features by interpreting the final results. The dataset we are given contains 84 (49%) divorced women/men and 86 (51%) women/men in happy marriage, each with 54 responses which includes detailed information about their marriage. In our research, we would mainly apply the correlation-based feature selection models on the whole dataset (`divorce.csv`).

## 2 Introduction and Problem Definition

The goal of this project is to apply different machine learning methods we've learned in class and develop different models to predict and shrink `divorce` model. This project consists of two parts: predict the divorce and choose the most significant features. For each part, we divide the dataset into training data set and test data set with the percentage of 70% and 30%, and apply the classification methods to predict the results.

1. In order to predict the `divorce` based on all features, we mainly apply classification methods and build the following models: Logistic Regression; Bayes' theorem-based methods including Linear Discriminant Analysis (LDA) and Naive Bayes; non-parametric method: K-nearest Neighborhoods (KNN); shrinkage methods including Lasso and Ridge Regression; tree-based method including Decision Tree, Bagging Tree and Random Forest; Support Vector Method including Support Vector Classifier and Support Vector Machine; Artificial Neural Network.
2. To identify significant features, we mainly apply Lasso Regression, Tree Pruning, Subset Selection and testify the efficiencies of each model.

## 3 Classification and Testing with All Features

### 3.1 Logistic Regression Method

We try to use multiple logistic regression to predict the binary response based on 54 quantitative predictors, and generalize the model as follows:

$$\log\left(\frac{p(X)}{1 - p(X)}\right) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p \quad (1)$$

where  $X = (X_1, X_2, \dots, X_{54})$  are 54 predictors,  $p(X)$  is the probability of divorce and  $1 - p(X)$  is the probability of a happy marriage. According to R, we get the test error rate is 7.69%.

However, due to the fact that the data are well-separated and maximum likelihood estimates are not unique, fitted probabilities numerically 0 or 1 occurred, we cannot conclude that the multiple logistic regression model as a good fit for our data.

### 3.2 Bayes' Theorem-based Method

While logistic regression directly generate a model for `divorce`, Bayes' theorem indicates another way of modeling by using conditional distribution.

**Linear Discriminant Analysis** In LDA, it models the distribution of predictors  $X$  separately in each response class ( $Y$ ), and then apply Bayes' theorem into estimates for  $Pr(Y = k | X = x)$ :

$$Pr(Y = k | X = x) = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)} \quad (2)$$

LDA assumes each predictor is normally distributed, so each class has the same variance. In this method, library `MASS` is used, with test error rate of 1.92%.

**Naive Bayes** Naive Bayes has been one of the simplest but also most powerful algorithms for classification modeling, it is a form of discriminant analysis that assumes features are independent in each class. It is especially useful when  $p$  is large, hence may perform well in our situation. For this method, we install packages `caret` and using both `caret` and `e1071` library, then plug into R with `NBclassifier = naiveBayes(Class~., data=divorce.train)` to build a model with test MSE of 0%.

### 3.3 Non-parametric Method

**KNN** The KNN classifier first identifies the  $K$  points in the training data that are closest to  $X_0$ , represented by  $N_0$ . It then estimates the conditional probability for class  $j$  as the fraction

of points in  $N_0$  whose response values equal  $j$ :

$$Pr(Y = j | X = x) = \frac{1}{K} \sum_{i \in N_0} I(y_i = j) \quad (3)$$

The choice of  $K$  has a drastic effect on the KNN classifier obtained. As  $K$  grows, the model becomes less flexible and produces a decision boundary that is close to linear. This represents the bias variance trade-off. Hence, we choose cross-validation to determine the value of  $K$ .

By dividing the data into 5 folds, we can obtain test error rate of 1.92%, so this result suggests this model is a good fit for our data.

### 3.4 Shrinkage Method

**Lasso Regression & Ridge Regression** With too many predictors, fitting the full model without penalization will result in large prediction intervals, and least square regression estimator may not uniquely exist. As an alternative, we can fit a model containing all  $p$  predictors by shrinking some of the coefficient estimates towards zero. The first approach is ridge regression, with  $\beta^R$  minimize :

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2 \quad (4)$$

$\lambda \sum_{j=1}^p \beta_j^2$  is a shrinkage penalty. In the R code, we use ridge regression to contain all the predictors and get final test error rate 1.92% which is same as the result of KNN approach.

The Lasso coefficients,  $\hat{\beta}_\lambda^L$  minimizes the quantity:

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j| = RSS + \lambda \sum_{j=1}^p |\beta_j| \quad (5)$$

As lasso could force some coefficients exactly equal to zero, it performs variable selection. In our case, the Lasso method obtains the same test error rate, 1.92%, as the ridge regression in divorce data.

### 3.5 Tree-based Method

**Decision Tree & Tree Pruning** Decision Tree is another powerful tool to manage our data. We predict that each observation belongs to the most commonly occurring class of training observations in the region to which it belongs. The classification error rate for one region is the fraction of the training observations in that region that do not belong to the most common class  $E = 1 - \max_k (\hat{p}_{mk})$  by checking the efficiency of this model, we got a test error rate of 3.85%.




Figure 1: Part of Decision Tree

Tree Pruning operates as a special case of Tree Decision Methods, while there might be chances that tree could overfit the data. In Tree Pruning, it reduces the chances of overfitting the training dataset by building a larger tree and then prune it to select a subtree which leads to the lowest test error rate. Each value of there corresponds a subtree  $T \subset T_0$  such that  $\sum_{m=1}^M \sum_{x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$  is as small as possible. By building the Tree Pruning model in R, we also got a test error rate of 3.85%, which is exactly the same as Tree Decision.

**Bagging** To further reduce variance of the decision tree method discussed above, bagging method is applied to the training dataset. We consider a bagging tree with 500 decision trees in total, each split of the tree holds all 54 predictors, then apply majority vote to assign a final value for `Class`.

$$\hat{f}_{bag}(X) = \text{MajorityVote}\{\hat{f}^{yb}(X)\} \quad (6)$$

As prediction for `Class` is chosen as the most commonly occurring class among all 500 predictions, model variance is greatly reduced, indicating that we now have a more stable model. Test error rate for bagging method is 1.90%.

**Random Forest** To further look into the questionnaire, we could safely assume that the answer to some questions might be crucial for a happy marriage. For example, answers to question 7: “We are like two strangers who share the same environment at home rather than family” might strongly decide whether a couple would divorce or not, as such questions, in a sense, could tell if spouses are still in love.

Such attributes could be seen as strong predictors in the model, which means all predictors at each split of the tree are not necessary to consider. To improve accuracy of our prediction, random forest model is applied to decorrelate the bagging tree, with a subset of 7 predictors to consider at each split. Test error rate of random forest method for divorce mode is 1.90%.

### 3.6 SVM

Support Vector Classifier and Support Vector Machines are both intended for the binary classification setting in which there are two classes. They are both in the interest of greater robustness to individual observations, and better classification of most of the training observations. Both of these two methods can be represented as:

$$f(x) = \beta_0 + \sum_{i \in S} \alpha_i K(x, x_i) \quad (7)$$

Where  $S$  is the collection of indices of these support points and  $K$  is kernel.

When we wish to get a linear boundary, we use linear kernel and the resulting classifier is known as Support Vector Classifier. When we assume the boundary is non-linear, we use a polynomial kernel of degree  $d$ , or a radial kernel with parameter  $\lambda$ , the resulting classifier is called Support Vector Machine.

Applying these two methods to the divorce data sets, we also use a cost argument, which allows us to specify the cost of a violation to the margin. When the cost argument is small, then the margins will be wide and many support vectors will be on the margin or will violate the margin. We then use cross-validation to choose the cost argument and parameter  $d$  or  $\lambda$ .

In our case, the estimate test error rates of Support Vector Classifier and Support vector Machine with radial kernel are both 1.92%, while the estimate test error rate of SVM with polynomial kernel is 0.00% (there is randomness in the result).

### 3.7 ANN

As some attributes might be stronger, artificial neural network is considered to improve the model. For better interpretation and lower computation assumption, feed-forward neural network is chosen in our analysis. In our case, answers to each attribute are collected as signals in input layer, with weight and activation function. This improves our model in a way that relationship between attributes inputs and their outputs are non-linear, which is correspondent with the real marriage situation; also, neurons in hidden layer serves as a filter for inputs, giving us a more efficient model.

To testify the improvement for different neural network models, four models have been considered in total: models with hidden layers and neurons of (1,1), (1,2), (2,1), (2,2). Graphic for neural network model with 1 hidden layer and 1 neuron is shown below:




Figure 2: Neural Network with 1-hidden layer, 1-neuron

Test error rates for the models are 1.92%, 1.92%, 0%, 1.92% separately, which implies neural network with 2-hidden layers, 1-neuron might be the best. As neural network model generates an output for `Class` as a probability value between 0 and 1, cross-entropy could be used to evaluate the models. The lowest cross-entropy errors among four models is 0.024 by model with one hidden layer and two neurons (the four errors are closed to each other, though).

## 4 Feature Selection

### 4.1 Stepwise Selection

We tried to find the important features in logistic regression by stepwise selection. We start with full model. In each step, we consider whether the  $i$ th variable belongs to the current variable set, if it exists, then we calculate AIC of the model after deleting this variable, if it does not exist, we calculate AIC of the model after adding this variable into current model. Each step we would get  $P$  different AIC for  $P$  different models, we chose the model that results in the smallest AIC and add or delete the relevant variable. We continue these steps until getting a model that no matter what variable we delete or add, we cannot get a smaller AIC. The variables in the final model are the variables that we consider important.

For the divorce data set, we use the training data set to perform stepwise selection, the variables remain in the final model are `Atr3`, `Atr26`, `Atr40`, which we think are important to logistic regression. We then build a logistic regression model based on these features, resulting in a 0 estimate test error rate. In this case, the feature selection may be effective.

## 4.2 Lasso Regression

As mentioned before, Lasso could force some coefficients exactly equal to zero, it performs variable selection. In our case, we apply Lasso to the training data set, the variables that has non-zero coefficients are Atr3, Atr6, Atr15, Atr16, Atr17, Atr18, Atr19, Atr20, Atr26, Atr30, Atr33, Atr39, Atr40, Atr49. We then build a logistic regression model based on these features selected by Lasso using the training data set, and test it with test data set. The estimate test error rate we get is 5.77%, which is smaller than 7.69%, the test error rate we get by applying logistic regression based on all features.

## 4.3 Tree Methods

From the result of tree pruning, possible important features are Atr17 and Atr20. From the importance measure of random forest, Atr26, Atr19, Atr17, Atr40 and Atr20 may be important.




Figure 3: Tree Pruning Selection

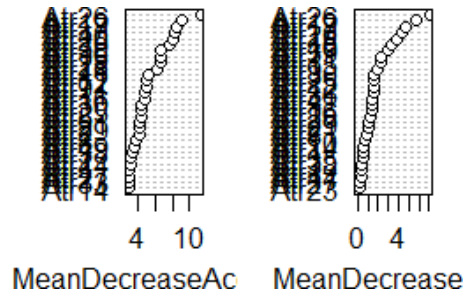


Figure 4: Random Forest Selection

## 5 Conclusion

In the analysis of `divorce` model, 12 machine learning methods are applied for variable prediction, with an average test error of 2.63%. To better compare test error rates between different models, a lollipop chart is attached below.




Figure 5: Test Error Rate for All Methods

As could be seen in the picture, both naive Bayes and neural network method have test error rate of 0. This might indicate they are the best two methods for model prediction; however, it can also be explained by the randomness of `set.seed()`, which gives us a well-separated testing dataset that perfectly corresponds with the classifier generated from the training set.

Also, as a result that `Class` values might be well separated as well, logistic regression method tends to have the largest test error rate among all methods.

Apart from these, test error rate for other methods are similar, which indicates that each of them could be seen as a strong and efficient method for building model `divorce` and predicting variable `Class`.

3 feature selection methods are applied to shrinkage the model into a more efficient one, each with a different subset of predictors while in general, some predictors occur more than once among all methods. For example, 3 (also the question: "When we need it, we can take our discussions with my spouse from the beginning and correct it"). This might imply that there exists multicollinearity among attributes. In our case, it tells that marriage situation is not determined by separated problems: when some problems occur, others appear as well, leading to a deterioration in marriage.

Based on the results above, we could safely conclude that questions in the questionnaire would greatly determine whether a happy marriage would exist, while some questions have stronger effect compared to others.



## 6 Appendix: R

```
1
2 #####
3 #1.predictthedivorcebasedonallthefeatures
4 #####
5
6 #readdata
7 divorce =read.csv('D:/divorce.csv')
8 #dividedataintotrainingssetandtesttest
9 #attach(divorce)
10 set.seed(2)
11 train =sample(1:nrow(divorce),0.7 *nrow(divorce))
12 Class =as.factor(Class)
13 divorce =cbind(divorce[, -ncol(divorce)],Class)
14 divorce.train = divorce[train,]
15 divorce.test = divorce[-train,]
16
17
18 #####
19 #Classificationmethod
20 #####
21
22 #logisticregressiontopredict:
23 glm.fit =glm(Class~.,data=divorce.train,
24             family=binomial,maxit=100)
25 contrasts(Class)
26 glm.pred =rep(0,nrow(divorce.test))
27 glm.probs =predict(glm.fit, newdata=divorce.test,
28                  type='response')
29 glm.pred[glm.probs > 0.5] = 1
30 mean(glm.pred!=Class[-train])
31 #fittedprobabilitiesnumerically0or1occurred
32
33 #LDA
34 library(MASS)
35 lda.fit = lda(Class~.,data=divorce.train)
36 lda.pred =predict(lda.fit, divorce.test)
37 lda.class= lda.pred$class
38 mean(lda.class!=Class[-train])
39 table(lda.class,Class[-train])
```

```

40
41 #naivebayes
42 library(caret)
43 NBclassifier = naiveBayes(Class~.,
44                           data=divorce.train)
45 NBpred =predict(NBclassifier,
46                newdata=divorce.test,
47                type='class')
48 mean(NBpred!=Class[-train])
49
50 #KNN
51 library(class)
52 #usecross-validationtochoosek
53 set.seed(2)
54 fold =sample(1:5,length(train),replace=T)
55 cv.errors=matrix(NA,5,floor(length(fold)/5),
56                 dimnames=list(NULL,paste(1:floor(length(fold)/5))))
57 for(i in 1:5){
58   for(j in 1:floor(length(fold)/5)){
59     knn.pred=knn(divorce.train[fold!=i,-ncol(divorce)],
60                 ,divorce.train[fold==i,-ncol(divorce)],
61                 Class[train][fold!=i],k=j)
62   cv.errors[i,j]=mean(knn.pred!=Class[train][fold==i])
63   }
64 }
65 mean.cv.errors=apply(cv.errors ,2,mean)
66 knn.k =which.min(cv.errors)
67 knn.pred=knn(divorce.train[,-ncol(divorce)],
68             divorce.test[,-ncol(divorce)],Class[train],k=knn.k)
69 knn.error.test =mean(knn.pred!=Class[-train])
70 knn.error.test
71
72
73 #####
74 #ModelSelectionmethod
75 #####
76 #Lasso
77 library(glmnet)
78 divorce =read.csv('D:/divorce.csv')
79 x=as.matrix(divorce)
80 y=x[,ncol(x)]

```

```

81 x=as.matrix(divorce)[,-ncol(x)]
82 cv.out=cv.glmnet(x[train,],y[train],alpha=1,
83                 family='binomial')
84 plot(cv.out)
85 bestlam=cv.out$lambda.min
86 lasso.mod=glmnet(x[train,],y[train],alpha=1,
87                 lambda=bestlam,family='binomial')
88 lasso.probs=predict(lasso.mod,s=bestlam,
89                   newx=x[-train,])
90 lasso.pred =rep(0,nrow(divorce.test))
91 lasso.pred[lasso.probs > 0.5] = 1
92 mean(lasso.pred!=Class[-train])
93
94 #ridge
95 cv.out=cv.glmnet(x[train,],y[train],alpha=0,
96                 family='binomial')
97 plot(cv.out)
98 bestlam=cv.out$lambda.min
99 lasso.mod=glmnet(x[train,],y[train],alpha=0,
100                 lambda=bestlam)
101 lasso.probs=predict(lasso.mod,s=bestlam,
102                   newx=x[-train,])
103 lasso.pred =rep(0,nrow(divorce.test))
104 lasso.pred[lasso.probs > 0.5] = 1
105 mean(lasso.pred!=Class[-train])
106
107
108 #####
109 #treemethods
110 #####
111 #Decisontree
112 library(tree)
113 tree.divorce = tree(Class~.,divorce.train)
114 #summary(tree.divorce)
115 plot(tree.divorce)
116 text(tree.divorce,pretty=0)
117 tree.pred =predict(tree.divorce,
118                   divorce.test, type='class')
119 mean(tree.pred!=Class[-train])
120
121 #Treepruning

```

```

122 set.seed(2)
123 cv.divorce=cv.tree(tree.divorce,FUN=prune.misclass)
124 prune.divorce = prune.misclass(tree.divorce,
125     best=cv.divorce$size[which.min(cv.divorce$dev)])
126 tree.pred =predict(prune.divorce,
127     divorce.test, type='class')
128 mean(tree.pred!=Class[-train])
129 plot(prune.divorce)
130 text(prune.divorce,pretty=0)
131 #didnotcutnodes
132
133 #baggingandrandomforests;
134 library(randomForest)
135 set.seed(2)
136 bag.divorce<-randomForest(Class~.,
137     data= divorce.train, mtry =ncol(divorce)-1,
138     importance = TRUE)
139 #bag.divorce
140 pred.bagging<-predict(bag.divorce,
141     newdata = divorce.test, type ="class")
142 table(pred.bagging, divorce.test$Class)
143 ER.bagging<-mean(pred.bagging!=divorce.test$Class)
144 cat("Testerrorrateforbaggingmethodis",
145     ER.bagging*100,"%.")
146
147 set.seed(2)
148 rf.divorce=randomForest(Class~.,
149     data=divorce.train,
150     mtry=floor(sqrt(ncol(divorce)-1)),
151     importance=TRUE)
152 yhat.rf=predict(rf.divorce,
153     newdata=divorce.test,type='class')
154 mean(yhat.rf!=Class[-train])
155 importance(rf.divorce)
156 #toseewhichpredictorsaremoreimportant.
157 varImpPlot(rf.divorce)
158
159 #####
160 #SVM
161 #####
162 #Supportvectorclassifier

```

```

163 library(e1071)
164 set.seed(2)
165 tune.out=tune(svm,Class~.,data=divorce.train,
166 kernel="linear",
167 ranges=list(cost=c(0.0001,0.001, 0.01, 0.1, 1,5,10,100)))
168 #summary(tune.out)
169 bestmod=tune.out$best.model
170 #summary(bestmod)
171 ypred=predict(bestmod,divorce.test)
172 mean(ypred!=Class[-train])
173
174 #Supportvectormachine
175 set.seed(2)
176 tune.out=tune(svm, Class~.,
177 data=divorce.train, kernel="radial",
178 ranges=list(cost=c(0.1,1,10,100,1000),
179 gamma=c(0.01,0.1,0.5,1,2,3,4)))
180 #summary(tune.out)
181 ypred=predict(tune.out$best.model,divorce.test)
182 mean(ypred!=Class[-train])
183
184 set.seed(2)
185 tune.out=tune(svm, Class~.,data=divorce.train,
186 kernel="polynomial",
187 ranges=list(cost=c(0.1,1,10,100,1000),
188 degree=c(2,3,4,5)))
189 #summary(tune.out)
190 ypred=predict(tune.out$best.model,divorce.test)
191 mean(ypred!=Class[-train])
192
193
194 #####
195 #neuralnetwork
196 #####
197 #1-hiddenlayer,1-neuron:
198 library(neuralnet)
199 library(tibble)
200 library(ggplot2)
201 set.seed(2)
202 model.nn1<-neuralnet(Class~.,
203 data= divorce.train, linear.output = FALSE,

```

```

204 err.fct = "ce", likelihood = TRUE)
205 # "linear.output&err.fct" for classification;
206 # "likelihood" to see AIC&BIC;
207 plot(model.nn1, rep = "best")
208
209 nn1.CE <- model.nn1$result.matrix[1, 1]
210 nn1.AIC <- model.nn1$result.matrix[4, 1]
211 nn1.BIC <- model.nn1$result.matrix[5, 1]
212 # see training error & AIC&BIC.
213
214 paste("For NN_model1, Cross-Entropy Error is",
215       round(nn1.CE, 3), ", AIC is", round(nn1.AIC, 3),
216       ", BIC is", round(nn1.BIC, 3), ".")
217
218 pred.nn1 <- compute(model.nn1, divorce.test)
219 # predict using model.nn1.
220 pred.result <- data.frame(ActualValue = divorce.test$Class,
221                          prediction = pred.nn1$net.result[, 2])
222 # see the prediction results in a data frame.
223 nn1.prediction <- round(pred.nn1$net.result[, 2], 0)
224 table(Actual = divorce.test$Class, Prediction = nn1.
225       prediction)
226 # see actual values and prediction results in a table.
227 ER.nn1 <- mean(divorce.test$Class != nn1.prediction)
228 # test error rate.
229 cat("Test error rate for neural network method
230     with 1-hidden layer is",
231     ER.nn1*100, "%.")
232
233 # compare with 1-hidden layer, 2-neurons;
234 # 2-hidden layers, 1-neuron; 2-hidden layers, 2-neurons;
235 # 1-layer, 2-neurons:
236 set.seed(2)
237 model.nn2 <- neuralnet(Class ~ ., data = divorce.train,
238                       linear.output = FALSE, err.fct = "ce",
239                       likelihood = TRUE, hidden = c(1, 2))
240 nn2.CE <- model.nn2$result.matrix[1, 1]
241 nn2.AIC <- model.nn2$result.matrix[4, 1]
242 nn2.BIC <- model.nn2$result.matrix[5, 1]
243 paste("For NN_model2, Cross-Entropy Error is",
244       round(nn2.CE, 3), ", AIC is", round(nn2.AIC, 3),

```

```

244         ",BICis",round(nn2.BIC, 3),".")
245
246 pred.nn2<-compute(model.nn2, divorce.test)
247 nn2.prediction<-round(pred.nn2$net.result[, 2], 0)
248 ER.nn2<-mean(divorce.test$Class!=nn2.prediction)
249 cat("Testerrorrateforneuralnetworkmethod
250 with1-hiddeenlayer,2-neuronsis", ER.nn2      *100,"%.")
251
252 #2-layers,1-neuron:
253 set.seed(2)
254 model.nn3<-neuralnet(Class~.,data= divorce.train,
255     linear.output = FALSE, err.fct ="ce",
256     likelihood = TRUE, hidden =c(2,1))
257 nn3.CE<-model.nn3$result.matrix[1, 1]
258 nn3.AIC<-model.nn3$result.matrix[4, 1]
259 nn3.BIC<-model.nn3$result.matrix[5, 1]
260 paste("ForNN_model3,Cross-EntropyErroris",
261     round(nn3.CE, 3),",AICis",round(nn3.AIC, 3),
262     ",BICis",round(nn3.BIC, 3),".")
263
264 pred.nn3<-compute(model.nn3, divorce.test)
265 nn3.prediction<-round(pred.nn3$net.result[, 2], 0)
266 ER.nn3<-mean(divorce.test$Class!=nn3.prediction)
267 cat("Testerrorrateforneuralnetworkmethod
268 with2-hiddeenlayers,1-neuronis", ER.nn3      *100,"%.")
269
270 #2-layers,2-neurons:
271 set.seed(2)
272 model.nn4<-neuralnet(Class~.,
273     data= divorce.train, linear.output = FALSE,
274     err.fct ="ce", likelihood = TRUE, hidden =c(2,2))
275 nn4.CE<-model.nn4$result.matrix[1, 1]
276 nn4.AIC<-model.nn4$result.matrix[4, 1]
277 nn4.BIC<-model.nn4$result.matrix[5, 1]
278 paste("ForNN_model4,Cross-EntropyErroris",
279     round(nn4.CE, 3),",AICis",round(nn4.AIC, 3),
280     ",BICis",round(nn4.BIC, 3),".")
281
282 pred.nn4<-compute(model.nn4, divorce.test)
283 nn4.prediction<-round(pred.nn4$net.result[, 2], 0)
284 ER.nn4<-mean(divorce.test$Class!=nn4.prediction)

```

```

285 cat("Testerrorrateforneuralnetworkmethod
286 with2-hiddeenlayers,2-neuronsis", ER.nn4      *100,"%.")
287
288 #seeinalollipopplot:
289 library(ggplot2)
290
291 data<-data.frame(
292   x=c("LogisticRegression","LDA","NaiveBayes",
293       "KNN","Lasso","RidgeRegression",
294       "DecisionTree","Bagging","RandomForest",
295       "SVC","SVM","ANN"),
296   y=c(7.69, 1.92, 0, 1.92, 1.92, 1.92,
297       3.85, 1.92, 1.92, 1.92, 1.92, 1.92)
298 )
299
300 ggplot(data, aes(x=x, y=y, label = paste0(y,"%")) +
301   geom_segment( aes(x=x, xend=x, y=0, yend=y),
302                 color="skyblue") +
303   geom_point( color="darkcyan", size=12, alpha=0.6) +
304   theme_light() +
305   coord_flip() +
306   geom_text(color = "white", size = 3)+
307   theme(
308     panel.grid.major.y = element_blank(),
309     panel.border = element_blank(),
310     axis.ticks.y = element_blank()
311   ) +
312   xlab("NameofMethods") +
313   ylab("TestErrorRate(%)")
314
315 #####
316 #Choosingimportantfeatures
317 #####
318
319 #####
320 #stepwiseselection
321 #####
322
323 mink=0
324 stepselect =function(data){
325   np =dim(data)

```



```

326 n = np[1]
327 p = np[2]-1
328 xn =names(data)[1:p]
329 #data=cbind(1,data)
330 glm.fit =glm(Class~.,data=data,
331             family=binomial)
332 hhb =glm.fit$coefficients
333 #loglik=logLik(glm.fit)
334 #AIC=-2 *loglik+2 *(p+1)
335 aic = AIC(glm.fit)
336 A = 1:p
337 MAIC = aic
338 mAIC = aic
339 flag =rep(FALSE,p)
340 #markwhetherthevariableisdeleted,false=deleted
341 repeat{
342     if(length(A) < p ){
343         B = (1:p)[flag]
344     }else{
345         B = NULL
346     }
347     AB =c(A,B)
348     AICm =rep(0,p)
349 #recordAICafteraddingorremovingthekthvariable
350     ff = 1
351     for(k in AB){
352         if(!flag[k]){
353             tA = A[-ff]
354             formula=paste('Class~1',
355                           paste(paste('Atr',tA,sep='')
356                                ,collapse ='+'),sep='+')
357             glm.fitk =glm(formula,
358                           data=data,family=binomial)
359             #loglikk=logLik(glm.fitk)
360             #AICk=-2 *loglikk+2 *(length(A)-1)
361             AICk = AIC(glm.fitk)
362             AICm[k] = AICk
363             ff = ff+1
364             if( AICk < mAIC ){
365 #ifAICisbecomingsmaller,recordcurrentk,
366                 #variableset,AIC

```

```

367         mtA = tA
368         mAIC = AICk
369
370     }
371 }else{
372     tA<-c(A,k)
373     formula=paste('Class~1',
374                   paste(paste('Atr',tA,sep=''),
375                         collapse ='+'),sep='+')
376     glm.fitk =glm(formula,
377                  data=data,family=binomial)
378     #loglikk=logLik(glm.fitk)
379     #AICk=-2 *loglikk+2 *(length(A)-1)
380     AICk = AIC(glm.fitk)
381     AICm[k] = AICk
382     if( AICk < mAIC ){
383         mink = k
384         mtA = tA
385         mAIC = AICk
386     }
387 }
388 }
389 if(mAIC >= MAIC)break
390 #ifAICnotdecreaseafteraddingorremovingvariable
391 if(mAIC<MAIC ){
392     flag[mink] =!flag[mink]
393     A = mtA
394     MAIC = mAIC#smallestAIC
395 }
396 }
397 formulaglm =paste('Class~1'
398                  ,paste(paste('Atr',A,sep=''),
399                          collapse ='+'),sep='+')
400 glm.fit =glm(formulaglm,data=data,
401              family=binomial)
402 b =glm.fit$coefficients
403 bm =rep(0,(p+1))
404 bm[c(1,A+1)]<-b
405 hhb = bm
406 glmopt =paste('Atr',A,sep='')
407 return(glmopt)

```

```

408 }
409
410 coefi2=stepselect(divorce.train)
411 #logisticregressiontopredict:
412 glm.fit =glm(Class~.,
413              data=divorce.train[,c(coefi2,'Class')],
414              family=binomial,maxit=100)
415 contrasts(Class)
416 glm.pred =rep(0,nrow(divorce.test))
417 glm.probs =predict(glm.fit,
418                  newdata=divorce.test[,c(coefi2,'Class')],
419                  type='response')
420 glm.pred[glm.probs > 0.5] = 1
421 mean(glm.pred!=Class[-train])
422
423
424 #####
425 #Lasso
426 #####
427 divorce =read.csv('D:/divorce.csv')
428 x=as.matrix(divorce.train)
429 y=x[,ncol(x)]
430 x=as.matrix(divorce.train[, -ncol(x)])
431 cv.out=cv.glmnet(x,y,alpha=1,family='binomial')
432 bestlam=cv.out$lambda.min
433 lasso.mod=glmnet(x,y,alpha=1,lambda=bestlam)
434 lasso.coef=coef(lasso.mod)[,1]
435 coefi =names(lasso.coef[lasso.coef!=0])[-1]
436 #testwith
437 #logisticregressiontopredict:
438 glm.fit =glm(Class~.,
439              data=divorce.train[,c(coefi,'Class')],
440              family=binomial,maxit=100)
441 contrasts(Class)
442 glm.pred =rep(0,nrow(divorce.test))
443 glm.probs =predict(glm.fit,
444                  newdata=divorce.test[,c(coefi,'Class')],
445                  type='response')
446 glm.pred[glm.probs > 0.5] = 1
447 mean(glm.pred!=Class[-train])

```